

Networking Network Services Project Work

OpenStack



Team members:

- > Rufat Isgandarov
- ➤ Yusif Hajizada
- > Vugar Hasanov

Introduction

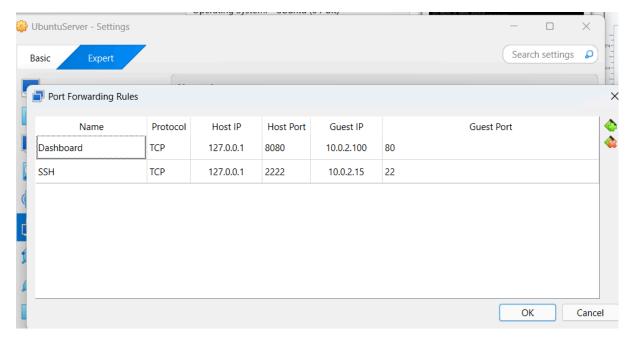
In this project, we wanted to learn how to install and run OpenStack using Kolla-Ansible in a virtual environment Oracle Virtualbox. OpenStack is a widely used open-source cloud computing platform that enables the creation and management of both public and private clouds. The objective of this work was to gain practical experience with the installation and configuration of OpenStack services, as well as to understand how its core components interact to provide Infrastructure-as-a-Service (laaS). So we decided to give it a try. We used a VirtualBox VM, installed Ubuntu, and followed the instructions to set everything up inside the VM. We used commands, config files, YAMLs. We followed structured steps including system preparation, network configuration, dependency installation, and the execution of Kolla-Ansible commands to build the OpenStack environment. We learned how different OpenStack services work together. Through this process, we became familiar with key services like Keystone, Nova, Neutron, Glance, and Horizon.

Main

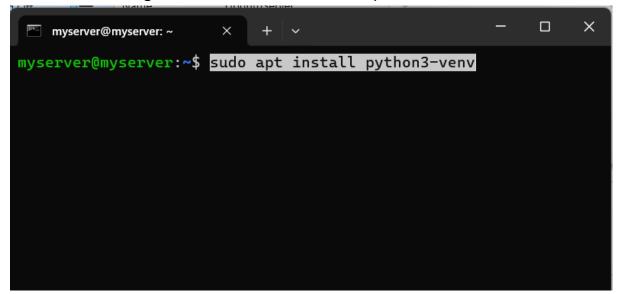
Install dependencies:

First of all, we specified Port Forwarding to connect VM with ssh from Guest machine terminal. We opened up VirtualBox's settings for our Ubuntu VM and went to the "Network" tab and configured it related to specification in lab description. This setup allows us to connect to the VM using SSH by typing the command <code>ssh ubuntu@localhost -p 2222</code> from our main computer. This made it easier for us to use the terminal without needing to open the VirtualBox window every time.

```
C:\Users\Rufat>ssh -p 2222 myserver@127.0.0.1 myserver@127.0.0.1's password:
```



Now we Installing the virtual environment dependencies.



Create a virtual environment and activate it:

Once we were in the VM, we created a new Python virtual environment using this command: python3 -m venv kolla/virtualenv This command created a folder where Python packages can be installed separately from the system-wide packages. Then we did source kolla/virtualenv/bin/activate so that python and pip now point into that venv.

```
myserver@myserver:~ × + v — — — ×

54 updates can be applied immediately.
To see these additional updates run: apt list --upgradable

Enable ESM Apps to receive additional future security updates.
See https://ubuntu.com/esm or run: sudo pro status

New release '24.04.2 LTS' available.
Run 'do-release-upgrade' to upgrade to it.

Last login: Sat Apr 26 09:27:21 2025

myserver@myserver:~$ python3 -m venv kolla/virtualenv
myserver@myserver:~$ source kolla/virtualenv/bin/activate
(virtualenv) myserver@myserver:~$
```

We noticed that the default pip version was outdated, so we decided to upgrade it to avoid issues during installation. We ran: *pip install -U pip* This updated pip to the latest version inside the virtual environment.

```
X
 myserver@myserver: ~
Last login: Sat Apr 26 09:27:21 2025
myserver@myserver:~$ python3 -m venv kolla/virtualenv
myserver@myserver:~$ source kolla/virtualenv/bin/activate
(virtualenv) myserver@myserver:~$ pip install -U pip
Requirement already satisfied: pip in ./kolla/virtualenv/lib/
python3.10/site-packages (22.0.2)
Collecting pip
  Downloading pip-25.0.1-py3-none-any.whl (1.8 MB)
                            - 1.8/1.8 MB 2.6 MB/s eta 0:00:00
Installing collected packages: pip
  Attempting uninstall: pip
    Found existing installation: pip 22.0.2
    Uninstalling pip-22.0.2:
      Successfully uninstalled pip-22.0.2
Successfully installed pip-25.0.1
(virtualenv) myserver@myserver:~$ |
```

Install Kolla-ansible. Install kolla-ansible and its dependencies using pip. With pip ready, we installed Kolla-Ansible directly from the official OpenDev repository. The command we used was:

pip install git+https://opendev.org/openstack/kolla-ansible@master
This downloaded and installed the latest Kolla-Ansible package and all
its dependencies.

Next, we created the folder /etc/kolla which will store all our configuration files. We did this using: This gave our user the right permissions to edit files inside /etc/kolla.

We copied the default example configuration files from the Kolla-Ansible installation directory to our newly created /etc/kolla folder.

Command used:

cp -r kolla/virtualenv/share/kolla-ansible/etc_examples/kolla/* /etc/kolla After this, we had globals.yml and passwords.yml ready to edit.

```
(virtualenv) myserver@myserver:~$ cp -r kolla/virtualenv/share/kolla-a
nsible/etc_examples/kolla/* /etc/kolla
(virtualenv) myserver@myserver:~$ ls kolla/
virtualenv
(virtualenv) myserver@myserver:~$ ls /etc/kolla/
globals.yml passwords.yml
(virtualenv) myserver@myserver:~$ |
```

To make the deployment process easier, we copied the all-in-one inventory file to our home directory:

cp kolla/virtualenv/share/kolla-ansible/ansible/inventory/all-in-one. This file is used by Ansible to know which hosts to install services on. Since we are using a single machine for everything, this file fits our needs perfectly.

```
(virtualenv) myserver@myserver:~$ cp kolla/virtualenv/share/kolla-ansi
ble/ansible/inventory/all-in-one .
(virtualenv) myserver@myserver:~$ ls
all-in-one kolla
(virtualenv) myserver@myserver:~$
```

Install Ansible Galaxy requirements. Kolla-Ansible uses some Ansible roles and plugins that are hosted on Ansible Galaxy. We installed them using: *kolla-ansible install-deps*

This downloaded all required roles.

```
all-in-one kolla
(virtualenv) myserver@myserver:~$ kolla-ansible install-deps
Installing Ansible Galaxy dependencies
Starting galaxy collection install process
Process install dependency map
Cloning into '/home/myserver/.ansible/tmp/ansible-local-1717ck73wq2r/t
mphhvl3vke/ansible-collection-kollaubxoluem'...
remote: Enumerating objects: 1230, done.
remote: Counting objects: 100% (477/477), done.
remote: Compressing objects: 100% (228/228), done.
remote: Total 1230 (delta 415), reused 249 (delta 249), pack-reused 75
```

Prepare initial configuration:

We ran *kolla-genpwd* to generate strong random passwords automatically and save them into */etc/kolla/passwords.yml*. This avoids

us having to set all the passwords manually. We checked the contents using cat to make sure everything was generated correctly.

Kolla passwords:

Passwords used in our deployment are stored in /etc/kolla/passwords.yml file.

```
(virtualenv) myserver@myserver:~$ kolla-genpwd
WARNING: Passwords file "/etc/kolla/passwords.yml" is world-readable.
The permissions will be changed.
(virtualenv) myserver@myserver:~$ |
```

```
(virtualenv) myserver@myserver:~$ cat /etc/kolla/passwords.yml
aodh_database_password: lGFx5DVbrfRbBItQIyh4hDLh9JZeZ3H4KQFkyVG9
aodh_keystone_password: sYuqUbWvvJMYOi9hRRCAFk8adtoCaXPRWVzyUrtm
barbican_crypto_key: wJUvm7MI-_9HV17vDsy8hSEyq9N0szJ_sTuLOXmqd94=
barbican_database_password: RbZwPMtlFK3rJlvzwWZRZRSQzfAhRf3eGNdNfKUA
barbican_keystone_password: emuSfIPNq2v3t3LKRUj6brmIosqUcEIQDjzXesnp
barbican_p11_password: D33DGCOe5gL8pgQRursBfdy36h3nZc46bbnW1LuA
bifrost_ssh_key:
    private_key: '-----BEGIN PRIVATE KEY-----
MIIJQQIBADANBgkqhkiG9w0BAQEFAASCCSswggknAgEAAoICAQDrOVK8dcKlvDnu
```


Kolla globals.yml globals.yml is the main configuration file for Kolla Ansible.

We opened the globals.yml file and set the following important settings:

- ➤ kolla base distro: "rocky"
- ➤ network_interface: "eth0"
- neutron_external_interface: "eth1"
- kolla_internal_vip_address: "10.0.0.50"

These settings define the Linux distribution we use, the interface for internal communication, and the interface for external access. We also set a virtual IP that services will use inside the network.

```
×
                                                                  myserver@myserver: ~
 GNU nano 6.2
                          /etc/kolla/globals.yml
# You can use this file to override _any_ variable throughout Kolla.
# Additional options can be found in the
# 'kolla-ansible/ansible/group_vars/all.yml' file. Default value of a>
# commented parameters are shown here, To override the default value >
# the parameter and change its value.
# Dummy variable to allow Ansible to accept this file.
workaround_ansible_issue_8743: yes
#####################
# Ansible options
####################
# This variable is used as the "filter" argument for the setup module
# instance, if one wants to remove/ignore all Neutron interface facts:
# kolla_ansible_setup_filter: "ansible_[!qt]*"
# By default, we do not provide a filter.
#kolla_ansible_setup_filter: "{{ omit }}"
# This variable is used as the "gather_subset" argument for the setup>
                          [ Read 866 lines ]
              ^O Write Out
                            ^W Where Is
                                           ^K Cut
^G Help
                                                           Execute
                                           ^U Paste
  Exit
              ^R Read File
                               Replace
                                                            Justify
```




Netplan configuration:

To make sure our VM had a fixed IP, we edited

/etc/netplan/01-netcfg.yaml and added static IP configuration for eth0. We included the IP address, netmask, gateway, and DNS server. After saving the file, we applied the settings using sudo netplan apply. This step ensured our server always keeps the same IP address after reboot.

```
myserver@myserver: ~
cat: /etc/netplan/50-cloud-init.yaml: Permission denied
myserver@myserver:~$ sudo cat /etc/netplan/*.yaml
# This file is generated from information provided by the datasource.
# to it will not persist across an instance reboot. To disable cloud-init'
# network configuration capabilities, write a file
 /etc/cloud/cloud.cfg.d/99-disable-network-config.cfg with the following:
# network: {config: disabled}
network:
    version: 2
    ethernets:
        enp0s3:
            dhcp4: true
        enp0s8:
            dhcp4: false
            addresses: [192.168.56.10/24]
            nameservers:
              addresses: [8.8.8.8, 1.1.1.1]
        enp0s9:
          dhcp4: false
          optional: true
        enp0s10:
          dhcp4: false
          optional: true
myserver@myserver:~$
```



```
X
  myserver@myserver: ~
                                        ×
  GNU nano 6.2
                                                     /etc/kolla/globals.yml
# Additionally, all vxlan/tunnel and storage network traffic will go over
# interface by default. This interface must contain an IP address.
# It is possible for hosts to have non-matching names of interfaces - thes>
# be set in an inventory file per host or per group or stored separately, >
           http://docs.ansible.com/ansible/latest/intro_inventory.html
# Yet another way to workaround the naming problem is to create a bond for>
# interface on all hosts and give the bond name here. Similar strategy can\gt
# followed for other types of interfaces.
network_interface: "enp0s8'
neutron_external_interface: "enp0s10"
# These can be adjusted for even more customization. The default is the sa>
# the 'network_interface'. These interfaces must contain an IP address.
#kolla_external_vip_interface: "{{    network_interface }}"
#api_interface: "{{ network_interface }}"
#tunnel_interface: "{{ network_interface }}"
#dns_interface: "{{ network_interface }}"
#octavia_network_interface: "{{ api_interface }}"
# Configure the address family (AF) per network.
# Valid options are [ ipv4, ipv6 ]
#network_address_family: "ipv4"
^G Help
                           ^O Write Out
                                                                                  ^K Cut
                                                                                                                  Execute
                                                       ^W Where Is
                           ^R
     Exit
                                Read File
                                                           Replace
                                                                                      Paste
                                                                                                                  Justify
                                                                                                                                        ×
  myserver@myserver: ~
myserver@myserver:~$ ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00
link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
inet 127.0.0.1/8 scope host lo
    valid_lft forever preferred_lft forever
inet6 ::1/128 scope host
    valid_lft forever preferred_lft forever

2: enp0s3: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
link/ether 08:00:27:90:22:15 brd ff:ff:ff:ff:
    inet 10.0.2.15/24 metric 100 brd 10.0.2.255 scope global dynamic enp0s3
    valid_lft 86132sec preferred_lft 86132sec
    inet6 fd00::a00:27ff:fe90:2215/64 scope global dynamic mngtmpaddr noprefixroute
    valid_lft 86133sec preferred_lft 14133sec
    inet6 fe80::a00:27ff:fe90:2215/64 scope link
    valid_lft forever preferred_lft forever
valid_lft forever preferred_lft forever
3: enp0s8: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
     link/ether 08:00:27:5b:63:87 brd ff:ff:ff:ff:ff:ff
inet 192.168.56.10/24 brd 192.168.56.255 scope global enp0s8
valid_lft forever preferred_lft forever
inet6 fe80::a00:27ff:fe5b:6387/64 scope link
4: enp0s9: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000 link/ether 08:00:27:ab:7b:d9 brd ff:ff:ff:ff
inet6 fe80::a00:27ff:feab:7bd9/64 scope link
         valid_lft forever preferred_lft forever
5: enp0s10: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000 link/ether 08:00:27:ed:5b:0e brd ff:ff:ff:ff:ff inet6 fe80::a00:27ff:feed:5b0e/64 scope link valid_lft forever preferred_lft forever
myserver@myserver:~$
```

Deployment

After configuration is set, we can proceed to the deployment phase. First we need to setup basic host-level dependencies, like docker. We prepared the server for deployment using this command:

kolla-ansible -i ./all-in-one bootstrap-servers

This step installed Docker and all other necessary packages.

We are doing pre-deployment checks for hosts:

Before starting the deployment, we wanted to make sure everything was ready. We used this command:

kolla-ansible -i ./all-in-one prechecks

This command checked network settings, disk space, user permissions, and other important things.

Now we were ready to install OpenStack services. We started the deployment with: *kolla-ansible -i ./all-in-one deploy*This took some time. It installed all the services such as Keystone, Glance, Nova, Neutron, and more

```
(virtualenv) myserver@myserver:~$ kolla-ansible deploy -i ./all-in-one |
                                  X
myserver@myserver: ~
skipping: no hosts matched
skipping: no hosts matched
[WARNING]: Could not match supplied host pattern, ignoring: enable_zun_True
skipping: no hosts matched
skipping: no hosts matched
skipping: no hosts matched
: ok=396 changed=258 unreachable=0 failed=0 skipped=236 rescued=
(virtualenv) myserver@myserver:~$
```

Deployment successful.

Now, To interact with OpenStack from the command line, we installed the OpenStack client (OpenStack CLI):

```
hon-keystoneclient, osc-lib, python-openstackclient
Attempting uninstall: urllib3
Found existing installation: urllib3 2.4.0
Uninstalling urllib3-2.4.0:
Successfully uninstalled urllib3-2.4.0
Attempting uninstall: tzdata
Found existing installation: tzdata 2025.2
Uninstalling tzdata-2025.2:
Successfully uninstalled tzdata-2025.2
Attempting uninstall: pyparsing
Found existing installation: pyparsing 3.2.3
Uninstalling pyparsing-3.2.3:
Successfully uninstalled pyparsing-3.2.3
Attempting uninstall: PrettyTable
Found existing installation: prettytable 3.16.0
Uninstalling prettytable-3.16.0:
Successfully uninstalled prettytable-3.16.0
Attempting uninstall: packaging
Found existing installation: packaging 25.0
Uninstalling packaging-25.0:
Successfully uninstalled prettytable-3.14.0
Uninstalling packaging-25.0:
Successfully uninstallic cryptography
Found existing installation: cryptography 44.0.2
Uninstalling cryptography-44.0.2
Successfully installed PrettyTable-3.14.0 cryptography-43.0.3 decorator-5.1.1 dogpile.cache-1.3.4 js
onpatch-1.33 jsonpointer-3.0.0 keystoneauth1-5.10.0 msgpack-1.1.0 openstacksdk-4.5.0 os-service-type
s-1.7.0 osc-lib-3.2.0 oslo.serialization-5.7.0 packaging-24.2 platformdirs-4.3.6 pyparsing-3.2.1 pyt
hon-cinderclient-9.7.0 python-keystoneolient-5.6.0 python-openstackclient-8.0.0 requestsexceptions-1
4.0 typing-extensions-4.12.2 tzdata-2025.1 urllib3-1.26.20
(virtualenv) myserver@myserver:-$
```

OpenStack requires a **clouds.yaml file** where credentials for the admin user are set. We ran kolla-ansible post-deploy to finish the setup. This created the admin-openrc.sh file with all the OpenStack environment variables. We also got the clouds.yaml file for using the OpenStack CLI. We copied the config to the right place with:

mkdir -p ~/.config/openstack && cp /etc/kolla/clouds.yaml ~/.config/openstack/

We ran this command to create a demo project, network, subnet, router, image, and a basic flavor: kolla/virtualenv/share/kolla-ansible/init-runonce This helped us quickly get a working setup to test instances Running init script to create demo resources:

```
(virtualenv) myserver@myserver:~$ ~/kolla/virtualenv/share/kolla-ansible/init-runonce
Checking for locally available cirros image.
None found, downloading cirros image (version 0.6.2).
% Total % Received % Xferd Average Speed Time Time Time Current
Dload Upload Total Spent Left Speed
                                     Θ
                                              Θ
                                                                    Θ --:--:--
                                                                                      0:00:01 -
100 20.4M 100 20.4M
Creating glance image.
                                              0 1481k
                                                                    0 0:00:14 0:00:14 --:-- 2344k
                                     Θ
 | Field
                                c8fc807773e5354afe61636071771906
   checksum
   container_format
                                bare
                                2025-04-27T09:45:32Z
   disk_format
                                qcow2
/v2/images/ba203f98-d0e6-4786-9fad-1706af7ae94d/file
   file
   id
                                ba203f98-d0e6-4786-9fad-1706af7ae94d
   min_disk
   min_ram
                                cirros
   name
   owner
                                53ece024cec34a91a28b32b7e359246a
                                os_hash_algo='sha512', os_hash_value='1103b92ce8ad966e41235a4de260deb791ff571
670c0342666c8582fbb9caefe6af07ebb11d34f44f8414b609b29c1bdf1d72ffa6faa39c88e87
   properties
                                21d09847952b', os_hidden='False', os_type='linux', owner_specified.openstack.md5='',
                                owner_specified.openstack.object='images/cirros',
owner_specified.openstack.sha256='', stores='file'
   protected
                                False
                                /v2/schemas/image
21430272
   schema
   size
   status
                                active
   tags
   updated_at
                                2025-04-27T09:45:32Z
   virtual_size
visibility
                                117440512
                                public
```

Configuring neutron.	
Field	Value
admin_state_up availability_zone_hints availability_zones	UP
created_at description	2025-04-27T09:45:36Z
distributed	False
enable_ndp_proxy	None null
external_gateway_info flavor_id	None
ha	False
id name	185b9edb-ab0e-4791-9849-04df53100b50 demo-router
project_id	53ece024cec34a91a28b32b7e359246a
revision_number	1
routes status	ACTIVE
tags	ACTIVE
tenant_id	53ece024cec34a91a28b32b7e359246a
updated_at	2025-04-27T09:45:36Z
+	+

+ Field	+
+	++ UP
availability_zone_hints	!
availability_zones	 2025-04-27T09:45:40Z
description	į į
dns_domain id	None 8ee8a750-20aa-4e56-91d8-6efa7e72fa06
ipv4_address_scope	None
ipv6_address_scope	None
is_default is_vlan_ging	None None
is_vlan_transparent	None
mtu name	1450 demo-net
port_security_enabled	True
project_id	53ece024cec34a91a28b32b7e359246a vxlan
provider:network_type provider:physical_network	VXLAN None
provider:segmentation_id	677
qos_policy_id revision_number	None
router:external	Internal
segments shared	None False
status	False ACTIVE
subnets	
tags updated_at	 2025-04-27T09:45:40Z
<u>+</u>	

Field	Value
allocation_pools	10.0.0.2-10.0.0.254
cidr	10.0.0.0/24
created_at	2025-04-27T09:45:43Z
description	
dns_nameservers	8.8.8.8
dns_publish_fixed_ip	None
enable_dhcp	True
gateway_ip	10.0.0.1
host_routes	
id	3175b850-f56e-4e9a-ab39-77e00636dcfe
ip_version	4
ipv6_address_mode	None
ipv6_ra_mode	None
name	demo-subnet
network_id	8ee8a750-20aa-4e56-91d8-6efa7e72fa06
project_id	53ece024cec34a91a28b32b7e359246a
revision_number	0
router:external	False
segment_id	None
service_types	
subnetpool_id	None
tags	
updated_at	2025-04-27T09:45:43Z
+	

Field	 Value
admin_state_up availability_zone_hints availability_zones	UP
created_at description	2025-04-27T09:45:51Z
dns_domain id	None 817fa8dd-a3db-4357-b5f5-45f5e25c8f66 None
ipv4_address_scope ipv6_address_scope is_default	None None False
is_vlan_qinq is_vlan_transparent	None
mtu name port_security_enabled	1500 public1 True
project_id provider:network_type	53ece024cec34a91a28b32b7e359246a flat
provider:physical_network provider:segmentation_id qos_policy_id	physnet1 None None
revision_number router:external	1 External
segments shared status	None False ACTIVE
subnets tags	
updated_at +	2025-04-27T09:45:52Z

+	Value
allocation_pools cidr created_at description dns_nameservers dns_publish_fixed_ip enable_dhcp	10.0.2.150-10.0.2.199 10.0.2.0/24 2025-04-27T09:45:54Z None False
gateway_ip host_routes id	10.0.2.1 2d559567-5f46-46ea-be10-eff96f086ba1
ip_version ipv6_address_mode ipv6_ra_mode	4 None None
name network_id project_id	public1-subnet 817fa8dd-a3db-4357-b5f5-45f5e25c8f66 53ece024cec34a91a28b32b7e359246a
revision_number router:external segment_id	θ True None
service_types subnetpool_id tags	None
updated_at	2025-04-27T09:45:54Z


+	
Field	Value
Field +	Value

```
Generating ssh key.
Generating public/private ecdsa key pair.
Your identification has been saved in /home/myserver/.ssh/id_ecdsa
Your public key has been saved in /home/myserver/.ssh/id_ecdsa.pub
The key fingerprint is:
SHA256:xIWQ9A8fvVMWppQMkY8mEN52scgSLklIRvI8cep5X68 myserver@myserver
The key's randomart image is:
+---[ECDSA 256]---+
 .0=.0.=+ .==..0
=.= +0*.0.=00 .
  = o +.X +oo o
  . o . +.=o..+
   o . Soo o
         Ε
+----[SHA256]----+
Configuring nova public key and quotas.
 Field
               | Value
  created_at
                None
  fingerprint
                 bf:27:5b:fc:c2:cd:a5:15:47:42:ec:a1:f0:e8:96:5f
  id
                 mykey
  is_deleted
                 None
  name
                 mykey
                 ssh
  type
  user_id
               4f478cff24f24a2e9eb76fb9c93332e5
```

Openstack administration


Access to Dashboard

We opened a browser and went to *http://192.168.56.200* which is the IP of our VM. The Horizon web interface loaded, and we logged in using the admin credentials from the *admin-openrc.sh*



```
(virtualenv) myserver@myserver:~$ cat /etc/kolla/admin-openrc.sh
# Ansible managed

# Clear any old environment that may conflict.
for key in $( set | awk '{FS="="} /^oS_/ {print $1}' ); do unset $key; done
export OS_PROJECT_DOMAIN_NAME='Default'
export OS_USER_DOMAIN_NAME='Default'
export OS_PROJECT_NAME='admin'
export OS_TENANT_NAME='admin'
export OS_USERNAME='admin'
export OS_PASSWORD='eYSXZP7VxZezTLuuBKPJJROdDEaQU4LtU72FMufO'
export OS_AUTH_URL='http://192.168.56.200:5000'
export OS_INTERFACE='internal'
export OS_ENDPOINT_TYPE='internalURL'
export OS_IDENTITY_API_VERSION='3'
export OS_REGION_NAME='RegionOne'
export OS_AUTH_PLUGIN='password'
```



Download the Cirros Image

We downloaded a test Cirros image *called cirros-0.6.2-x86_64-disk.img* and uploaded it. This made the image available to launch new virtual machines.


(virtualenv) myserver@myserver:~\$ wget https://download.cirros-cloud.net/0.6.2/cirros-0.6.2-x86_64-disk.img

The image Cirros via openstack CLI, is created.

```
rtualenv) myserver@myserver:~$ openstack image create "cirros"
-file cirros-0.6.2-x86_64-disk.img \
-disk-format qcow2 --container-format bare \
-nublic
   -public
Field
checksum
                                     c8fc807773e5354afe61636071771906
container_format
created_at
                                    bare
                                    2025-04-27T13:47:43Z
disk_format
file
id
                                    /v2/images/22ce4135-9591-4466-8fdd-7a4f0748d87c/file
22ce4135-9591-4466-8fdd-7a4f0748d87c
min_disk
min_ram
name
                                    cirros
                                     53ece024cec34a91a28b32b7e359246a
owner
                                    53ece024cec34491a2853207e339246a
os_hash_algo='sha512', os_hash_value='1103b92ce8ad966e41235a4de260deb791ff571670c0342666c
8582fbb9caefe6af07ebb11d34f444f8414b609b29c1bdf1d72ffa6faa39c88e8721d09847952b',
os_hidden='False', owner_specified.openstack.md5='',
owner_specified.openstack.object='images/cirros', owner_specified.openstack.sha256='',
stores='file'
False
''22/sebs_se_/inse_
properties
protected
                                    /v2/schemas/image
21430272
size
status
                                    active
tags
updated_at
                                    2025-04-27T13:47:45Z
virtual size
                                     117440512
 visibility
```

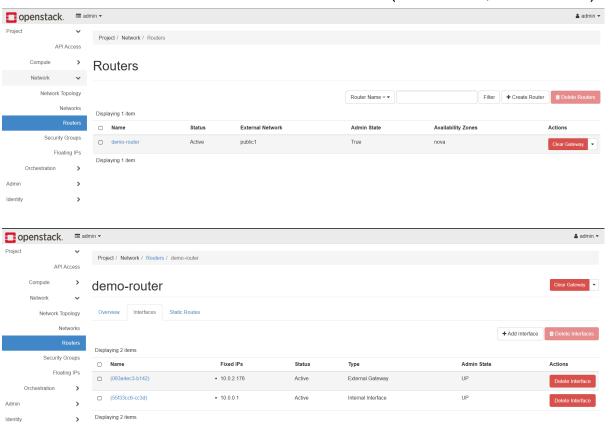
We created two networks:

- ➤ An internal private network called **demo-net** with IP range 10.0.0.0/24
- ➤ An external network called **public1** with IP range 10.0.2.0/24 (the one we will connect our router to).



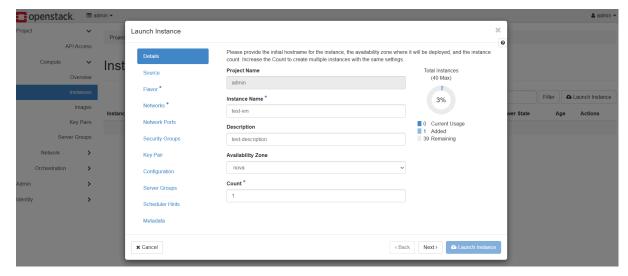
We created a router named router1, set its external gateway to public1, and added an interface connecting it to demo-net. This allowed traffic between the private network and external network. The router is connected to the External Network (public1, 10.0.2.0/24).

And it's also connected to the Internal Network (demo-net, 10.0.0.0/24).

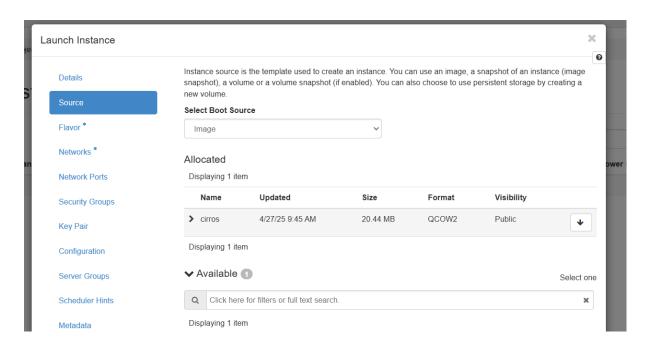


Launching the instance and testing the connectivity the external networks:

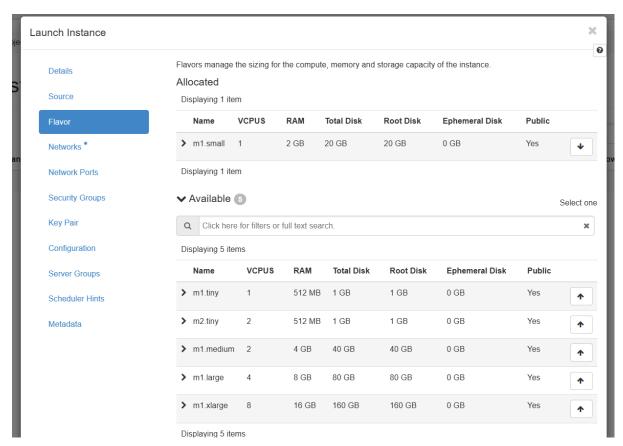
Using the Cirros image and demo network, we created a small virtual machine



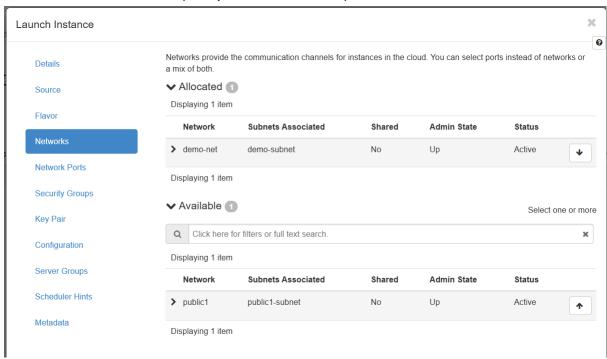
We selected the Cirros image that we uploaded earlier.



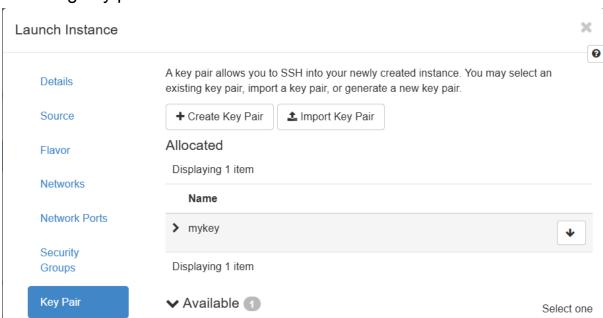
We selected Flavor: small one (m1.small)



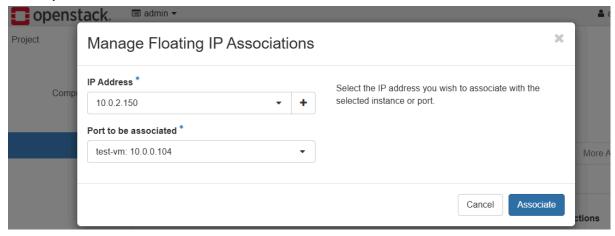
Attach it to demo-net (the private network).



Attaching key pair:



Associate the floating IP 10.0.2.150 to our instance (test-vm) in the OpenStack Dashboard:




The information: image, flavor and etc.

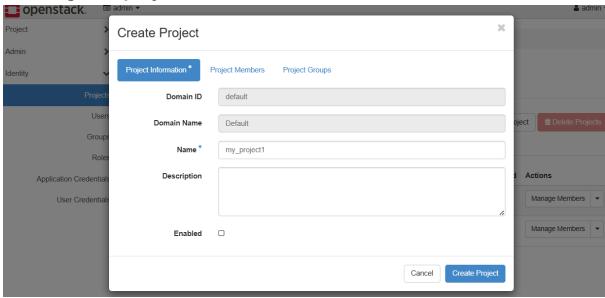
(virt	ualenv) myseı	rver@myse	erver:~	openst	tack i	lmage lis	st		
ID				Name	Status				
	++ 22ce4135-9591-4466-8fdd-7a4f0748d87c cirros active ba203f98-d0e6-4786-9fad-1706af7ae94d cirros active								
(virt	ualenv) mysei	rver@myse	erver:~9	openst	tack f	lavor l	ist		
ID	Name	RAM	Disk	Epheme	eral	VCPUs	Is Pu	blic	
1 2 3	m1.tiny m1.small m1.medium	512 2048 4096	1 20 40		0 0 0	1 1 2	True True True		
4 5 6	m1.large m1.xlarge m2.tiny	8192 16384 512	80 160 1		0 0 0	4 8 2	True True True		
(virtu	(virtualenv) myserver@myserver:~\$ openstack project list								
ID				Nar	ne	į			
	205ec5dd6c4d446c900f99e7b7e79163 service 53ece024cec34a91a28b32b7e359246a admin								
(virtu	(virtualenv) myserver@myserver:~\$ openstack keypair list								
Name	Name Fingerprint Type				+ Type	į			
	mykey bf:27:5b:fc:c2:cd:a5:15:47:42:ec:a1:f0:e8:96:5f ssh mykey2 a7:71:d0:a6:cc:07:2f:67:c7:75:70:d9:fa:61:86:8e ssh mykey2 a7:71:d0:a6:cc:07:2f:67:c7:75:70:d0:a6:cc:07:2f:67:d0:a6:cc:07:2f:67:a6:cc:07:d0:a6:cc:07:a6:cc:07:d0:a6:cc:07:cc:07:d0:a6:cc:07:cc:07:d0:								

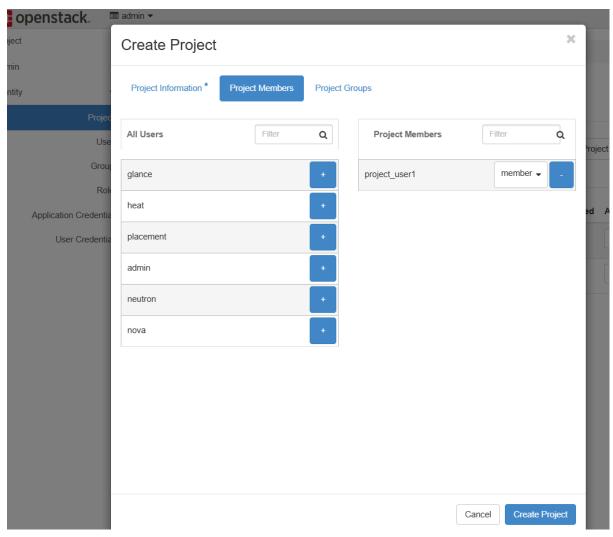
Openstack administration: manage user project

Creating new user:

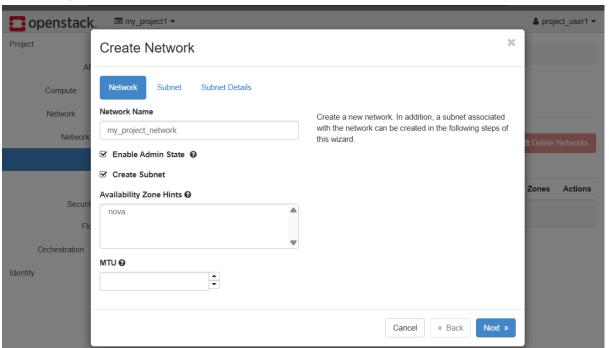


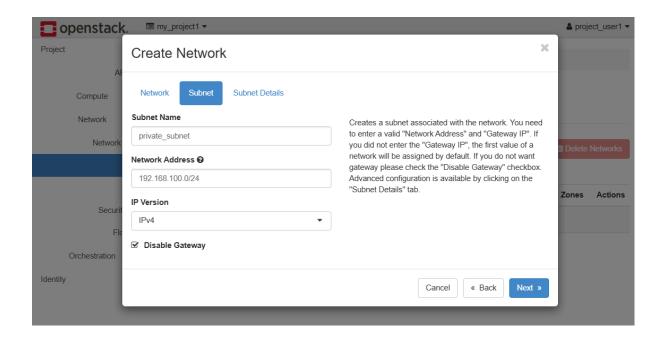
Creating new project:



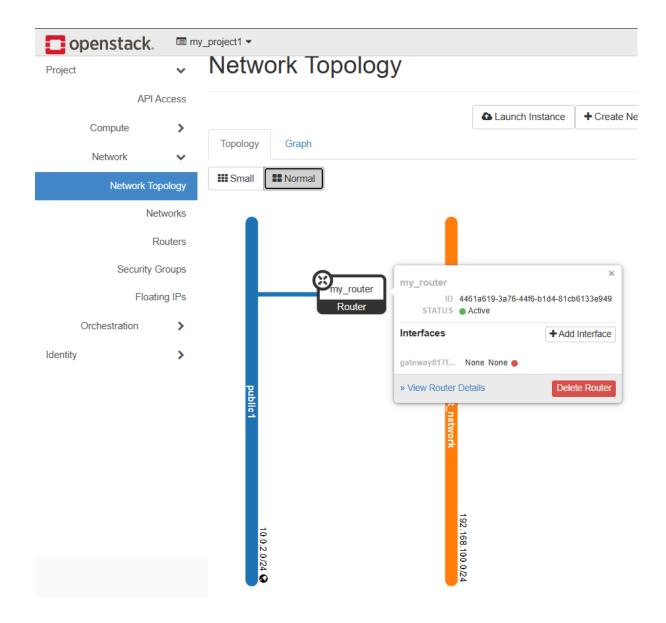


Creating network in new user:

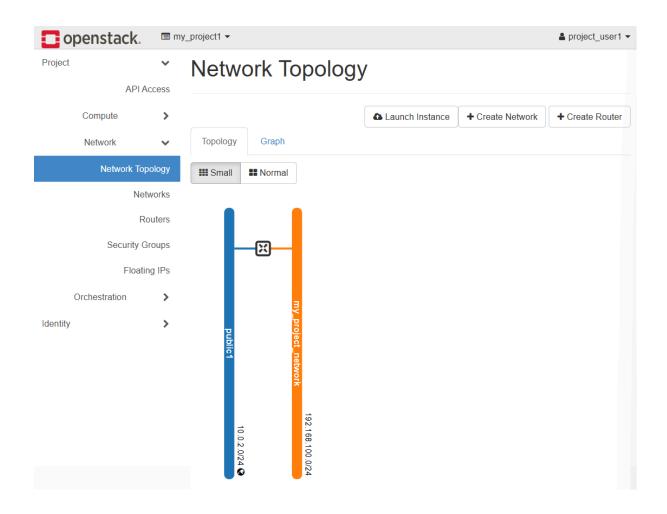




Displaying the Network topology and our router is connected to the « public network »



Displaying toplogy, we see that our private network is attached to the router:



Conclusion

After completion of the whole project, we feel way more confident with OpenStack. We gained valuable hands-on experience with OpenStack deployment using Kolla-Ansible. We now understand how to use Kolla-Ansible to deploy OpenStack on one single node VM, how to configure networks, launch instances, and manage cloud resources through both command-line tools and the Horizon dashboard. We also realized that YAML files. This project helped solidify our understanding of the architecture and functionality of OpenStack components. It also highlighted the importance of system preparation, proper configuration management, and troubleshooting skills when working with complex cloud platforms. This project served as a strong introduction to cloud infrastructure management and provided a solid foundation for more advanced OpenStack use cases in the future. In the end, everything worked, and we saw our VM running inside the cloud that we built ourselves.